

What can we learn from a **billion** agents?

Ayush Chopra
MIT

github.com/AgentTorch/AgentTorch

AI economist with 4 agents



Zheng et al (2021)

AI economist with 4 agents



Zheng et al (2021)

Smallville with 25 agents



Park et al (2023)

AI economist with 4 agents



Zheng et al (2021)

Smallville with 25 agents



Park et al (2023)

Minecraft with 1000 agents



Yang et al (2024)

AI economist with 4 agents



Zheng et al (2021)

Smallville with 25 agents



Park et al (2023)

Minecraft with 1000 agents



Yang et al (2024)

Deepmind framework for generative social simulations



Concordia

A library for generative social simulation

python 3.11 | 3.12 pyPI v1.7.0 pypi-test passing test-concordia passing test-examples passing

[Concordia Tech Report](#)

Vezhnevets et al (2023)

Imagine millions of autonomous agents

mimic global demographics + calibrated to reality



high-resolution agents + population-scale interactions



END THE
LOCKDOWN

VACCINATE
HIMSELF!

LOCKDOWN!

DON'T
BE A
LAB RA

COVID=
SCAM

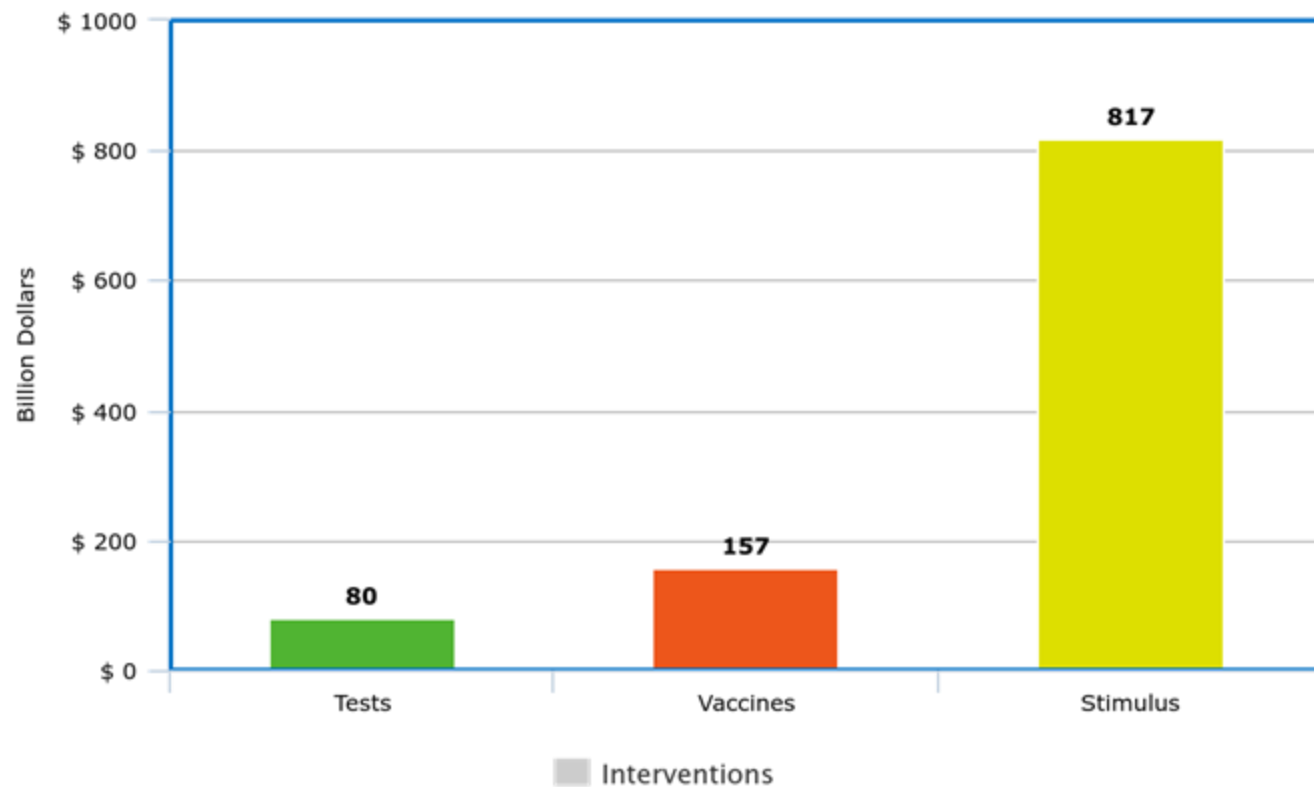
THIS SIGN
TESTED POSITIVE
FOR COVID-19

SAVE CALIFORNIA
RECALL GAVIN 2020
SIGN HERE

NA
VENTA

LT
LM
GAS
COVE
58

Trillion-Dollar Pandemic





"What kind of a test do you want?"

10% false-positive
3 day result

40% false positive
1 day result



Large Population Models

lpm.media.mit.edu/platform

model millions of interacting agents to capture collective dynamics



5 million citizens

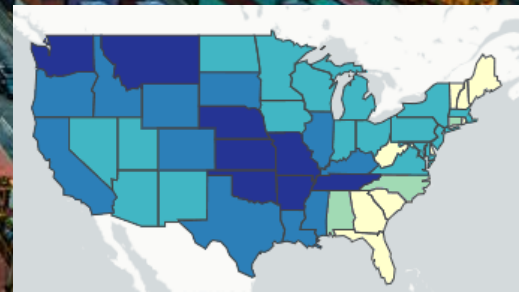


30 million households





\$ billion supply chains



Outline: What can we learn from a billion agents?

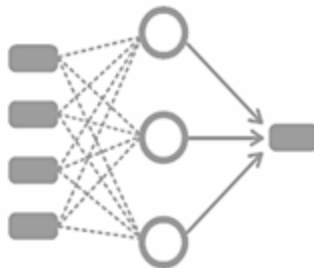
- Motivation and Impact
- Introduction
 - Large Population Models (LPMs)
 - The AgentTorch framework
- Case Study
 - Modeling disease spread over 8.4 million agents
- Conclusion

Challenges with agent-based simulations



Expressive agents

Adaptive
Heterogeneous



Dynamic interactions

Multi-scale
Stochastic



Synchronized analysis

Decentralized
Sensitive

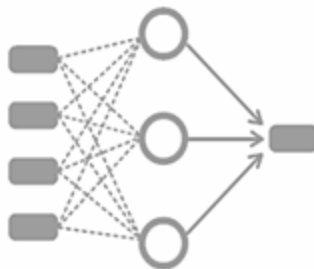
Large Population Models



Expressive agents

Adaptive
Heterogeneous

LLM Archetypes



Dynamic interactions

Multi-scale
Stochastic



Synchronized analysis

Decentralized
Sensitive

Large Population Models




Expressive agents

Adaptive
Heterogeneous

LLM Archetypes

AAMAS'23; arxiv'24



Dynamic interactions

Multi-scale
Stochastic

Differentiability

BMJ'21; AAMAS'23;
AAMAS'24 (Best paper runner-up)



Synchronized analysis

Decentralized
Sensitive

Large Population Models

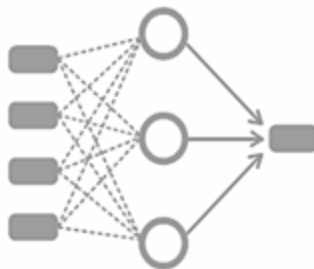


Expressive agents

Adaptive
Heterogeneous

LLM Archetypes

AAMAS'23; arxiv'24



Dynamic interactions

Multi-scale
Stochastic

Differentiability

BMJ'21; AAMAS'23;
AAMAS'24 (Best paper runner-up)



Synchronized analysis

Decentralized
Sensitive

Secure MPC

ICML-W'22 (Best paper);
AAMAS'24; Nature Medicine'25

Simulate an entire country on commodity hardware

Execute upto 300,000 interactions/sec and scale to 60 million agents/GPU

Method	Simulation	Calibration	Analysis
Conventional ABM*	50 hours	100,000 hours	5,000 hours
LPM	5 minutes	12 hours	10 seconds

Simulate an entire country on commodity hardware

Execute upto 300,000 interactions/sec and scale to 60 million agents/GPU

Method	Simulation	Calibration	Analysis
Conventional ABM*	50 hours	100,000 hours	5,000 hours
LPM	5 minutes	12 hours	10 seconds


+ 600x

+ 8300x

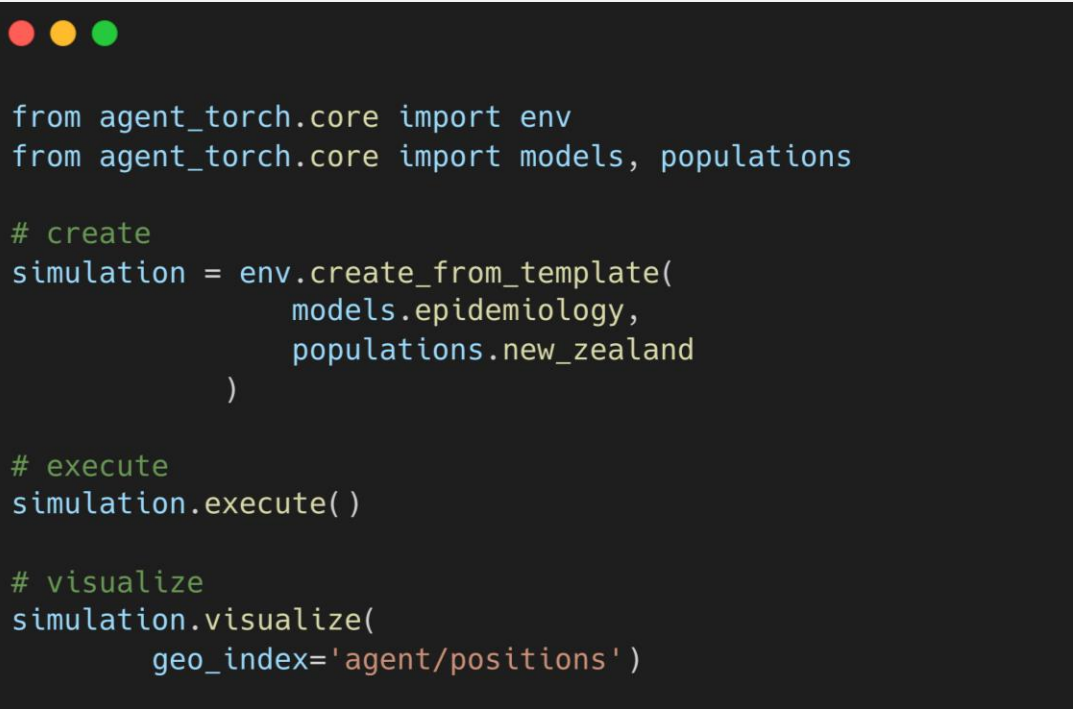
+ 5000x

github.com/AgentTorch/AgentTorch

`pip install agent-torch`



Build end-to-end workflows in 3 lines of code



```
from agent_torch.core import env
from agent_torch.core import models, populations

# create
simulation = env.create_from_template(
    models.epidemiology,
    populations.new_zealand
)

# execute
simulation.execute()

# visualize
simulation.visualize(
    geo_index='agent/positions')
```

Expressive templates

Custom populations

High-res visualizations

Build end-to-end workflows in 3 lines of code

```
from agent_torch.core import env
from agent_torch.core import models, populations

# create
simulation = env.create_from_template(
    models.epidemiology,
    populations.new_zealand
)


# execute
simulation.execute()

# visualize
simulation.visualize(
    geo_index='agent/positions')
```

Expressive templates

Custom populations

High-res visualizations



Comparison with other frameworks

Feature	AgentTorch	Concordia	Flame	Mesa
GPU Execution	✓	✓	✓	✗
Million-agent Populations	✓	✗	✓	⚠
Differentiable Environments	✓	✗	✗	✗
Mechanistic Environments	✓	✗	✓	✓
LLM Integration	✓	✓	✗	✗
Neural Composition	✓	⚠	✗	✗

Outline: What can we learn from a billion agents?

- Motivation and Impact
- Introduction
- Case Study:
 - Problem Formulation: modeling disease spread over 8.4 million agents
 - Part 1: Population Prompting via Archetypes
 - Part 2: Gradient-assisted simulation via Differentiable Environments
 - Part 3: Decentralized analysis via Secret Sharing
- Conclusion



END THE
LOCKDOWN

VACCINATE
HIMSELF!

LOCKDOWN!

DON'T
BE A
LAB RA

COVID=
SCAM

THIS SIGN
TESTED POSITIVE
FOR COVID-19

SAVE CALIFORNIA
RECALL GAVIN 2020
SIGN HERE

NO VEHICLE
ENTRY

58


New York City

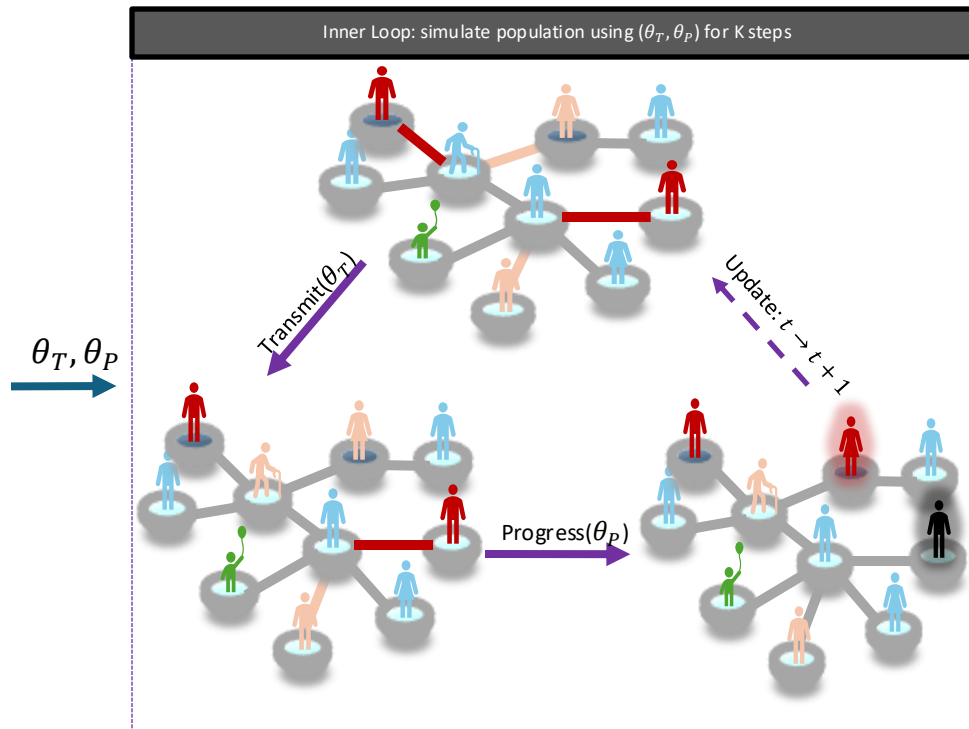
- 5 boroughs
- 8.4 million people
- 3.3 million households
- 200,000 small businesses

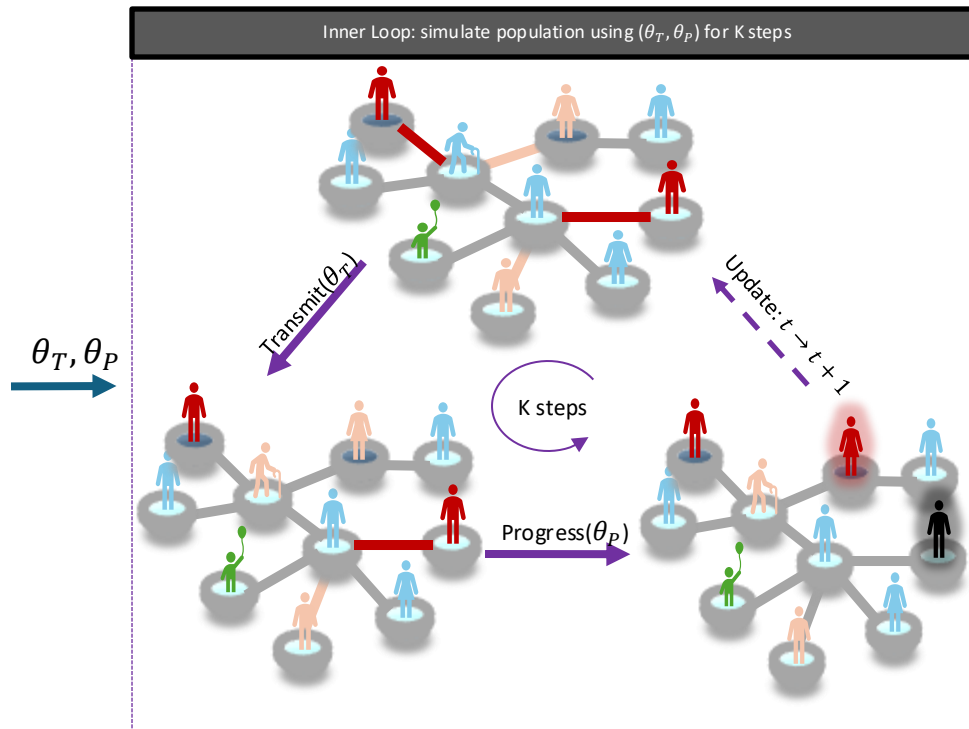


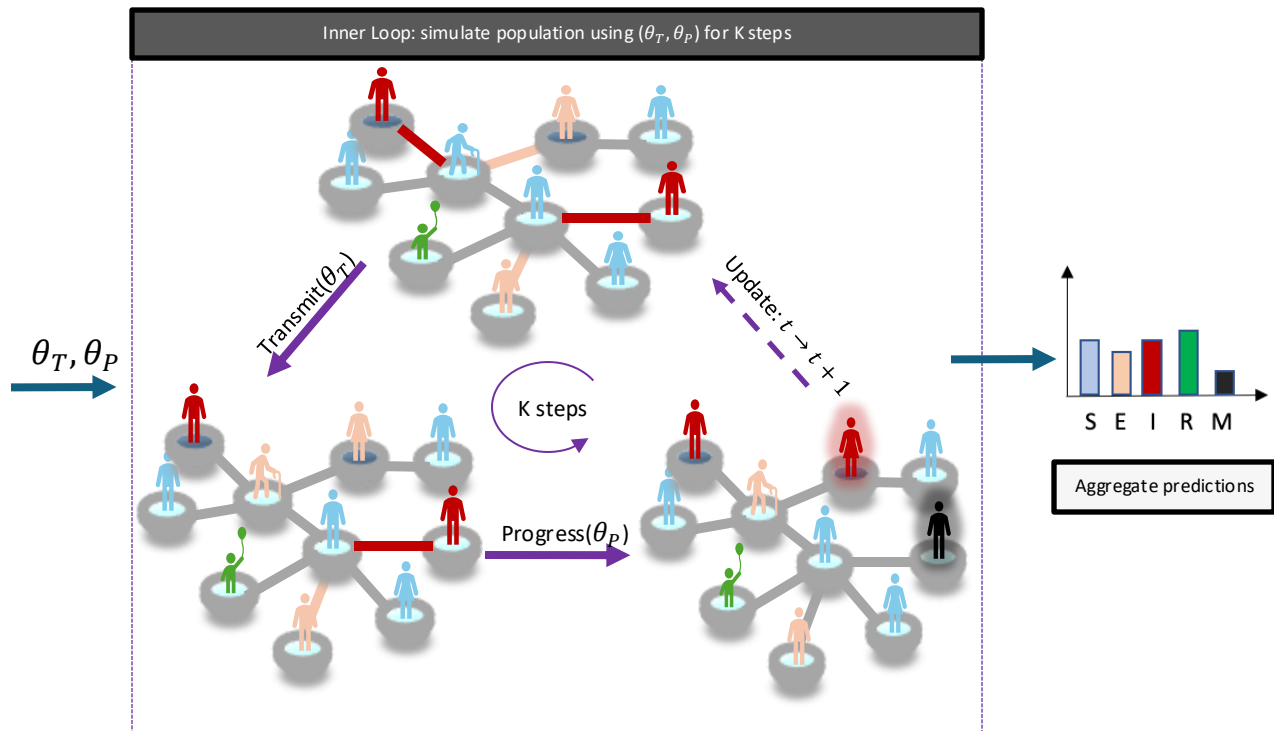
2020-2022

What if we give \$500 stimulus payments?









$$\mathbf{z}_i(t+1) = f\left(\mathbf{z}_i(t), \bigoplus_{j \in \mathcal{N}_i(t)} M_{ij}(t), \boldsymbol{\theta}, \ell(\cdot | \mathbf{z}_i(t))\right)$$

$$\mathbf{z}_i(t+1) = f \left(\mathbf{z}_i(t), \bigoplus_{j \in \mathcal{N}_i(t)} M_{ij}(t), \boldsymbol{\theta}, \ell(\cdot | \mathbf{z}_i(t)) \right)$$

Diagram illustrating the components of the state transition function f :

- New state of agent**: $\mathbf{z}_i(t+1)$
- Current state of agent**: $\mathbf{z}_i(t)$
- neighborhood influence**: $\bigoplus_{j \in \mathcal{N}_i(t)} M_{ij}(t)$
- environmental influence**: $\boldsymbol{\theta}$
- Individual agency**: $\ell(\cdot | \mathbf{z}_i(t))$

“Spread factor” of the variant
environmental influence

- Disease status
- Age, Income, Occupation

Current state of agent

$$\mathbf{z}_i(t+1) = f \left(\mathbf{z}_i(t), \bigoplus_{j \in \mathcal{N}_i(t)} M_{ij}(t), \boldsymbol{\theta}, \ell(\cdot | \mathbf{z}_i(t)) \right)$$

New state of agent

- New disease status
- Age, Income, Occupation

neighborhood influence

- Time since infection
- Vaccination status


Individual agency

Willingness to interact

Q1: model individual agency?

$$\mathbf{z}_i(t+1) = f \left(\mathbf{z}_i(t), \bigoplus_{j \in \mathcal{N}_i(t)} M_{ij}(t), \boldsymbol{\theta}, \ell(\cdot | \mathbf{z}_i(t)) \right)$$

- Varies with agent state
- Need to be sampled online



LLMs to prompt 8.4 millions agents in NYC?

Large Population Models



Expressive agents

Adaptive
Heterogeneous

LLM Archetypes

AAMAS'23; arxiv'24



Dynamic interactions

Multi-scale
Stochastic

Differentiability

BMJ'21; AAMAS'23;
AAMAS'24 (Best paper runner-up)



Data-driven analysis

Decentralized
Sensitive

Secure MPC

ICML-W'22 (Best paper);
AAMAS'24; Nature Medicine'25



Expressive agents

Adaptive
Heterogeneous

LLM Archetypes

Scale LLMs to prompt millions of agents

```
from agent_torch.core import Archetype, Behavior
from agent_torch.populations import NYC

# Create an object of Archetype class
# n_arch estimates a predictive posterior over outcomes
archetype = Archetype(n_arch=7)

# Create an object of Behavior class
work_behavior = Behavior(archetype=archetype.llm(prompt),
                        region=NYC)

will_work = work_behavior.sample()
```


Agent Prompt

User Prompt

You are a {gender} of age {age}, living in the {location} region. You work in {occupation } industry with monthly income of {income }.

The number of new cases in your neighborhood is {cases}, which is a {change}% change from the previous month. It has been {duration} months since the start of the pandemic.

This month, you have received a stimulus payment of {payment} to support your living expenses.

Given these factors, do you choose to isolate at home? (isolation behavior)

"There isn't enough information" and "It is unclear" are not acceptable answers. Give a "Yes" or "No" answer, followed by a period. Give one sentence explaining your choice.

#archetypes = age x gender x regions x occupation x income-levels

Specify prompt-specific archetypes



8.4 million agents in NYC



~2000 representative archetypes

prompt dimensions
data resolution (e.g. census)

Population Prompting via Archetypes






Prompt

•
•
•



Archetypes

Population Prompting via Archetypes




Prompt

Aggregate

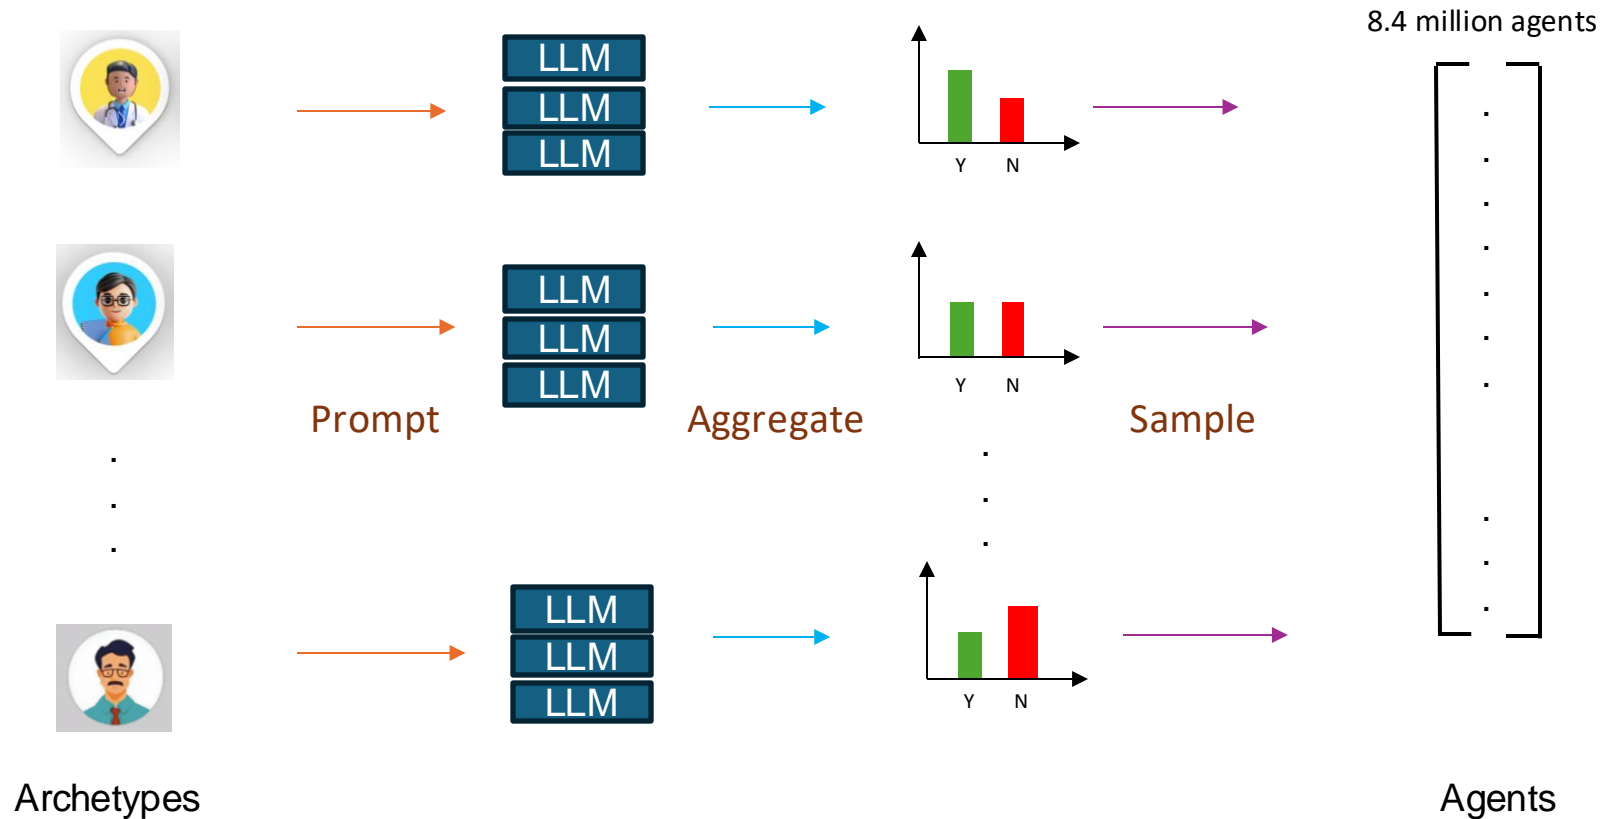
.
. .
. .

.
. .
. .



Archetypes

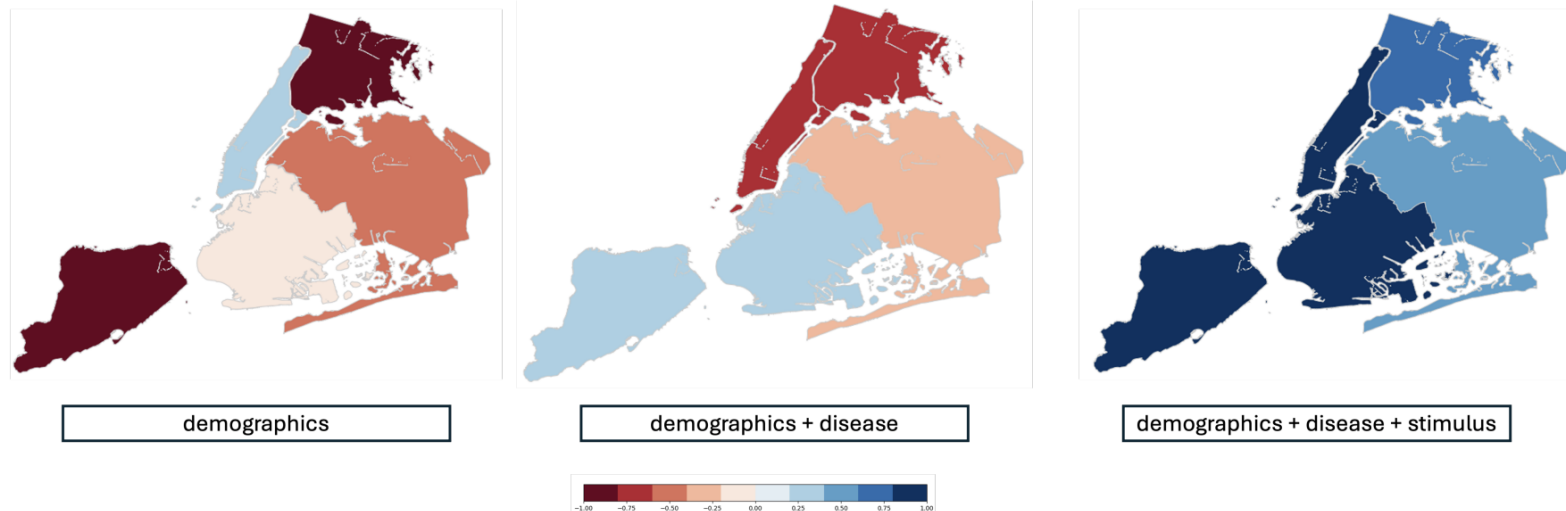
Population Prompting via Archetypes



High variance individual agents produce reliable census-level estimates

Prompt virtual population of 8.4 million agents

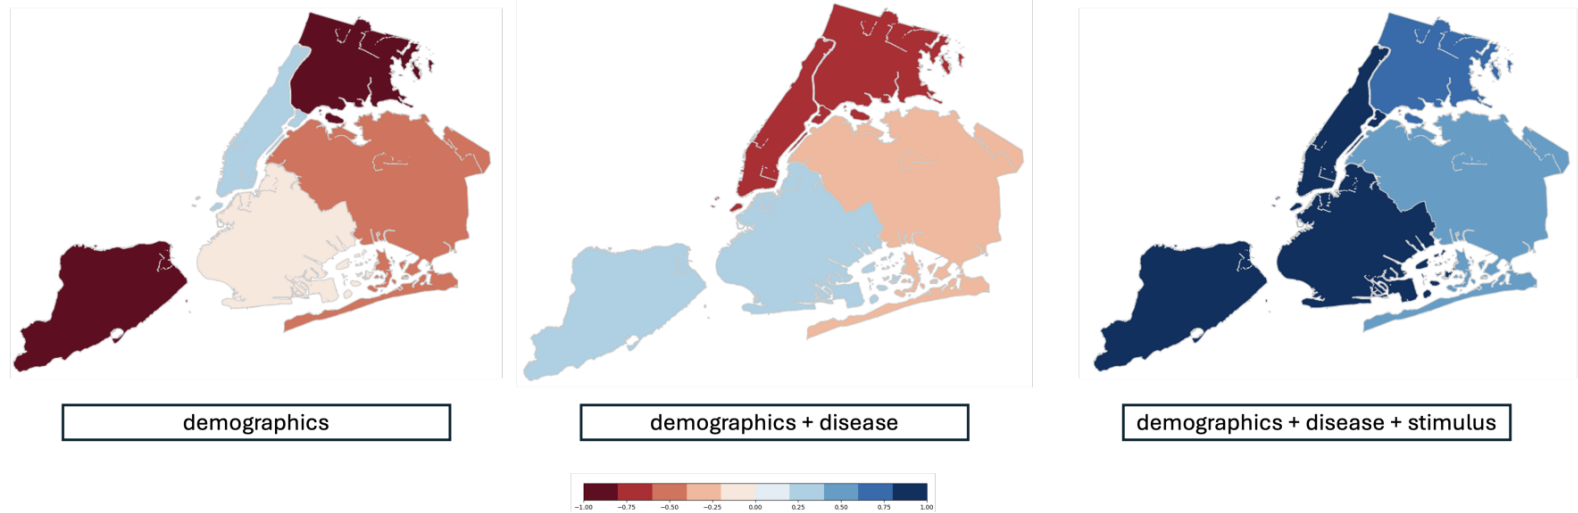
Measure correlation with labor force participation census



High variance individual agents produce reliable census-level estimates

Prompt virtual population of 8.4 million agents

Measure correlation with labor force participation census



Rebuild the census, in simulation!?

$$\mathbf{z}_i(t+1) = f \left(\mathbf{z}_i(t), \bigoplus_{j \in \mathcal{N}_i(t)} M_{ij}(t), \boldsymbol{\theta}, \ell(\cdot | \mathbf{z}_i(t)) \right)$$

Q2: simulate and calibrate efficiently?

Simulation involves repeated execution of f


Calibration involves estimating θ

$$\mathbf{z}_i(t+1) = f \left(\mathbf{z}_i(t), \bigoplus_{j \in \mathcal{N}_i(t)} M_{ij}(t), \theta, \ell(\cdot | \mathbf{z}_i(t)) \right)$$

Q2: simulate and calibrate efficiently?

Simulation involves repeated execution of f

Calibration involves estimating θ

$$\mathbf{z}_i(t+1) = f \left(\mathbf{z}_i(t), \bigoplus_{j \in \mathcal{N}_i(t)} M_{ij}(t), \theta, \ell(\cdot | \mathbf{z}_i(t)) \right)$$


- neural or mechanistic
- Stochastic and multiple “sub-steps”

f vectorized and differentiable?

Large Population Models




Expressive agents

Adaptive
Heterogeneous

LLM Archetypes

AAMAS'23; arxiv'24



Dynamic interactions

Multi-scale
Stochastic

Differentiability

BMJ'21; AAMAS'23;
AAMAS'24 (Best paper runner-up)




Data-driven analysis

Decentralized
Sensitive

Secure MPC

ICML-W'22 (Best paper);
AAMAS'24; Nature Medicine'25

Compose end-to-end with NNs



Dynamic interactions

Multi-scale
Stochastic

Differentiability



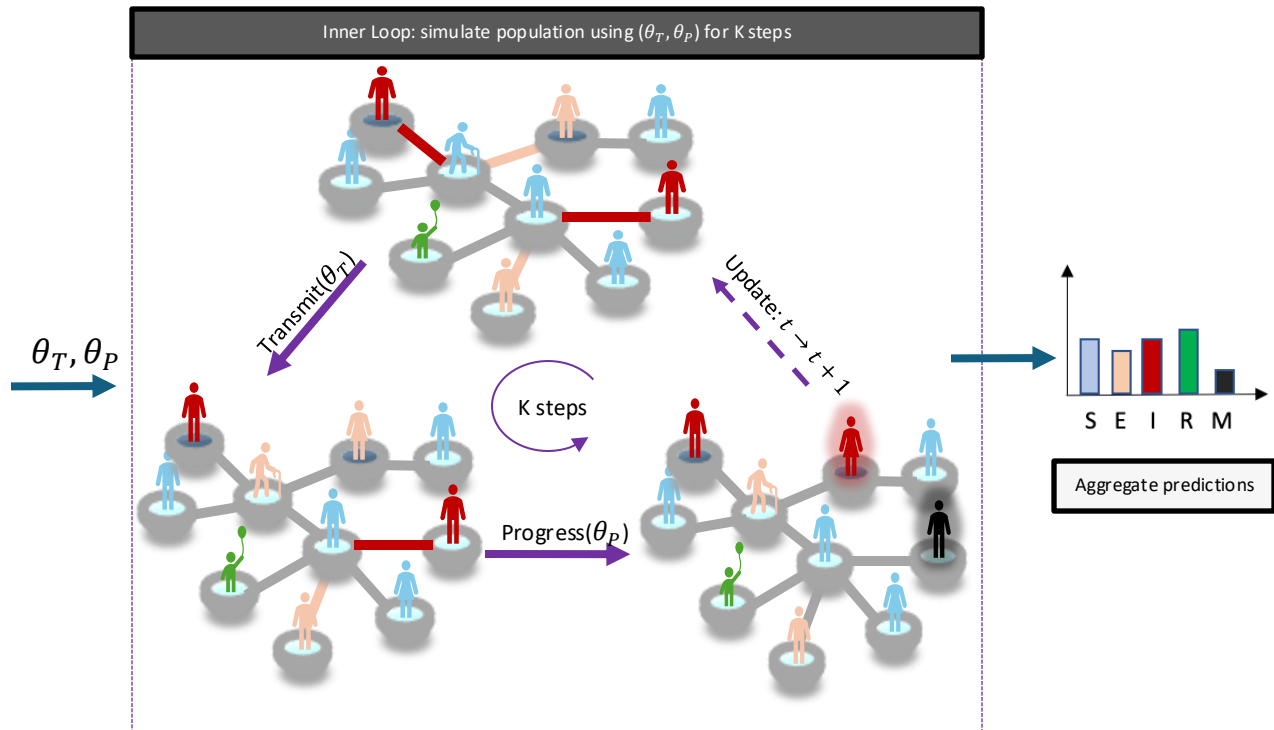
```
# agent_"torch" works seamlessly with the pytorch API
from torch.optim import SGD


nn = compose_nn()

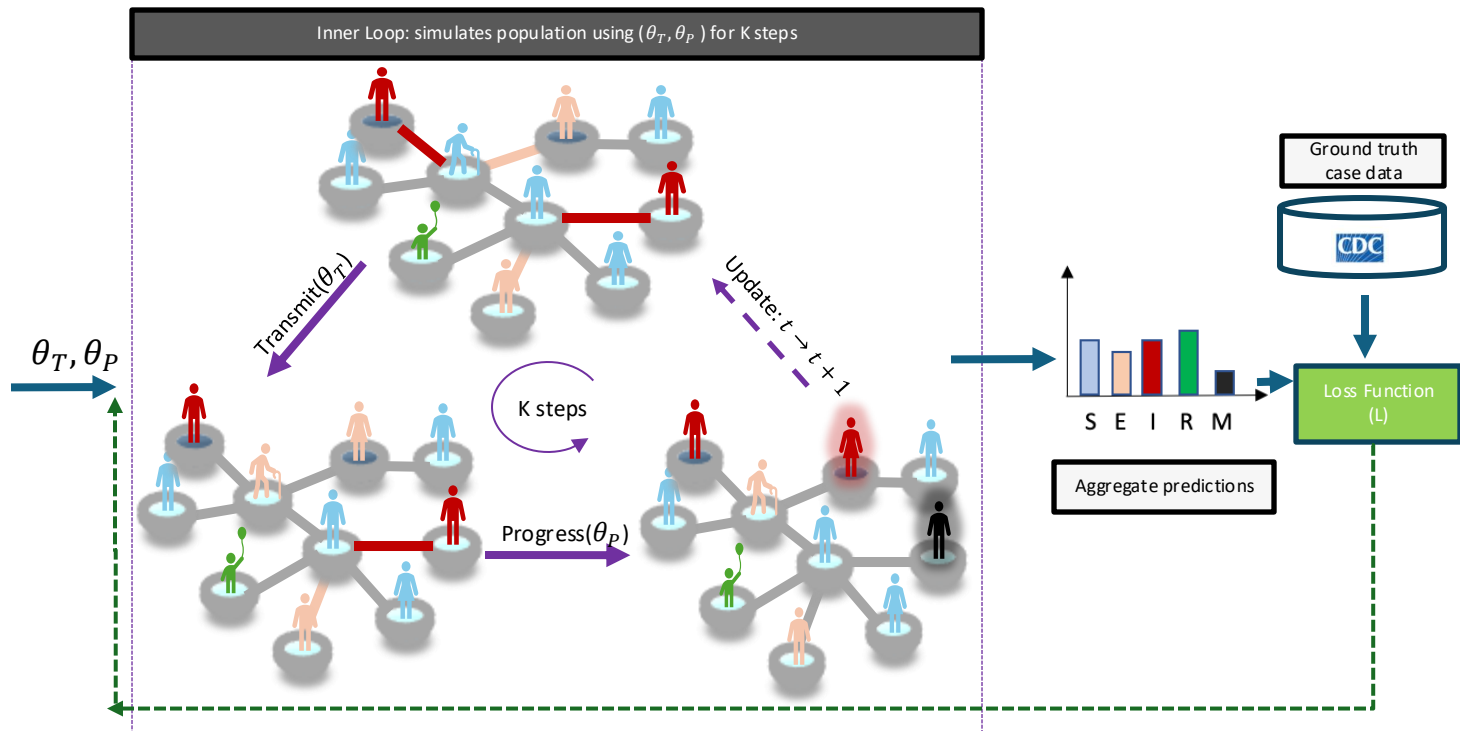
for i in range(n_epochs):
    parameters = compose_nn.forward()

    simulation.step(params=parameters)
    cases, employment = simulation.predict()

    simulation.optimize(SGD)
    simulation.reset()
```

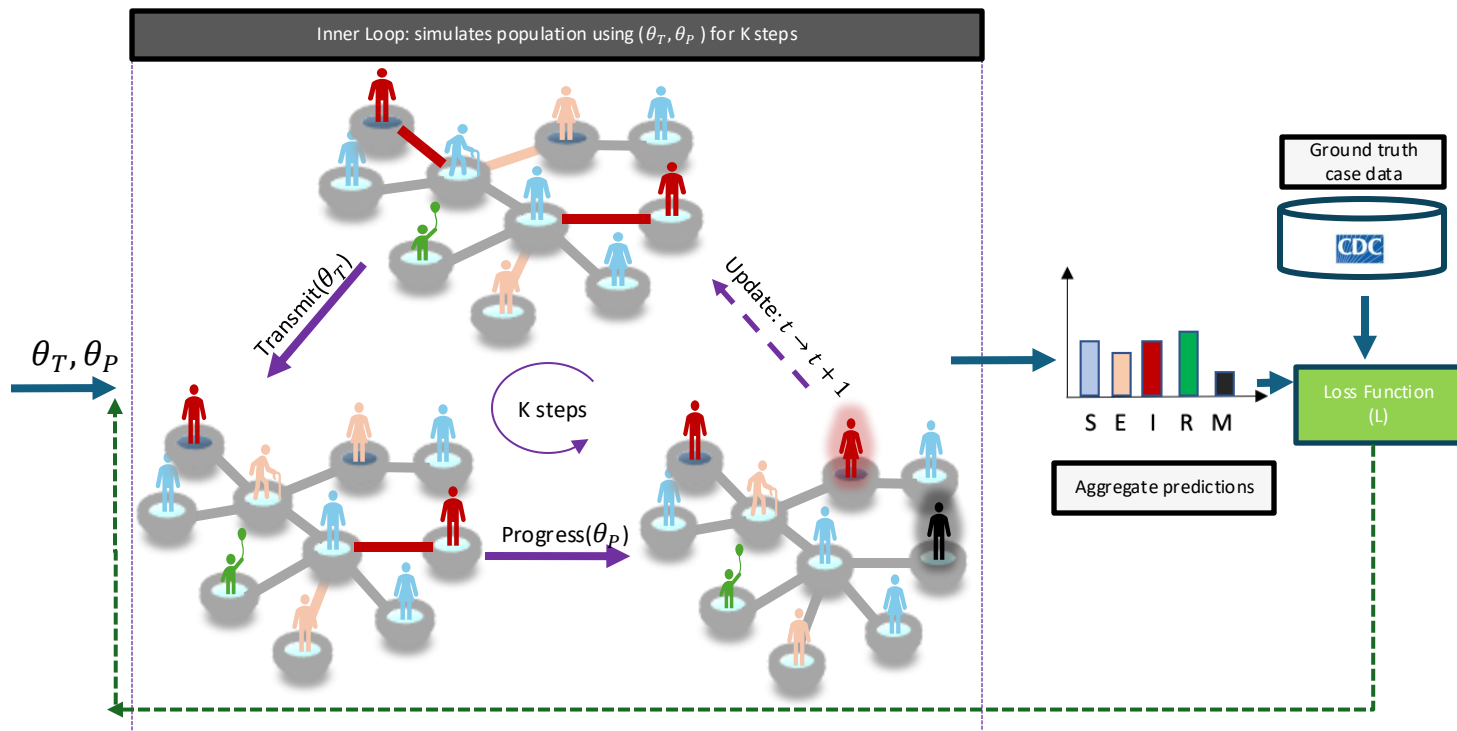






$$\theta_T = \theta_T - \alpha \frac{\partial \mathcal{L}(\hat{y}, y)}{\partial \theta_T}$$


$$\theta_P = \theta_P - \alpha \frac{\partial \mathcal{L}(\hat{y}, y)}{\partial \theta_P}$$

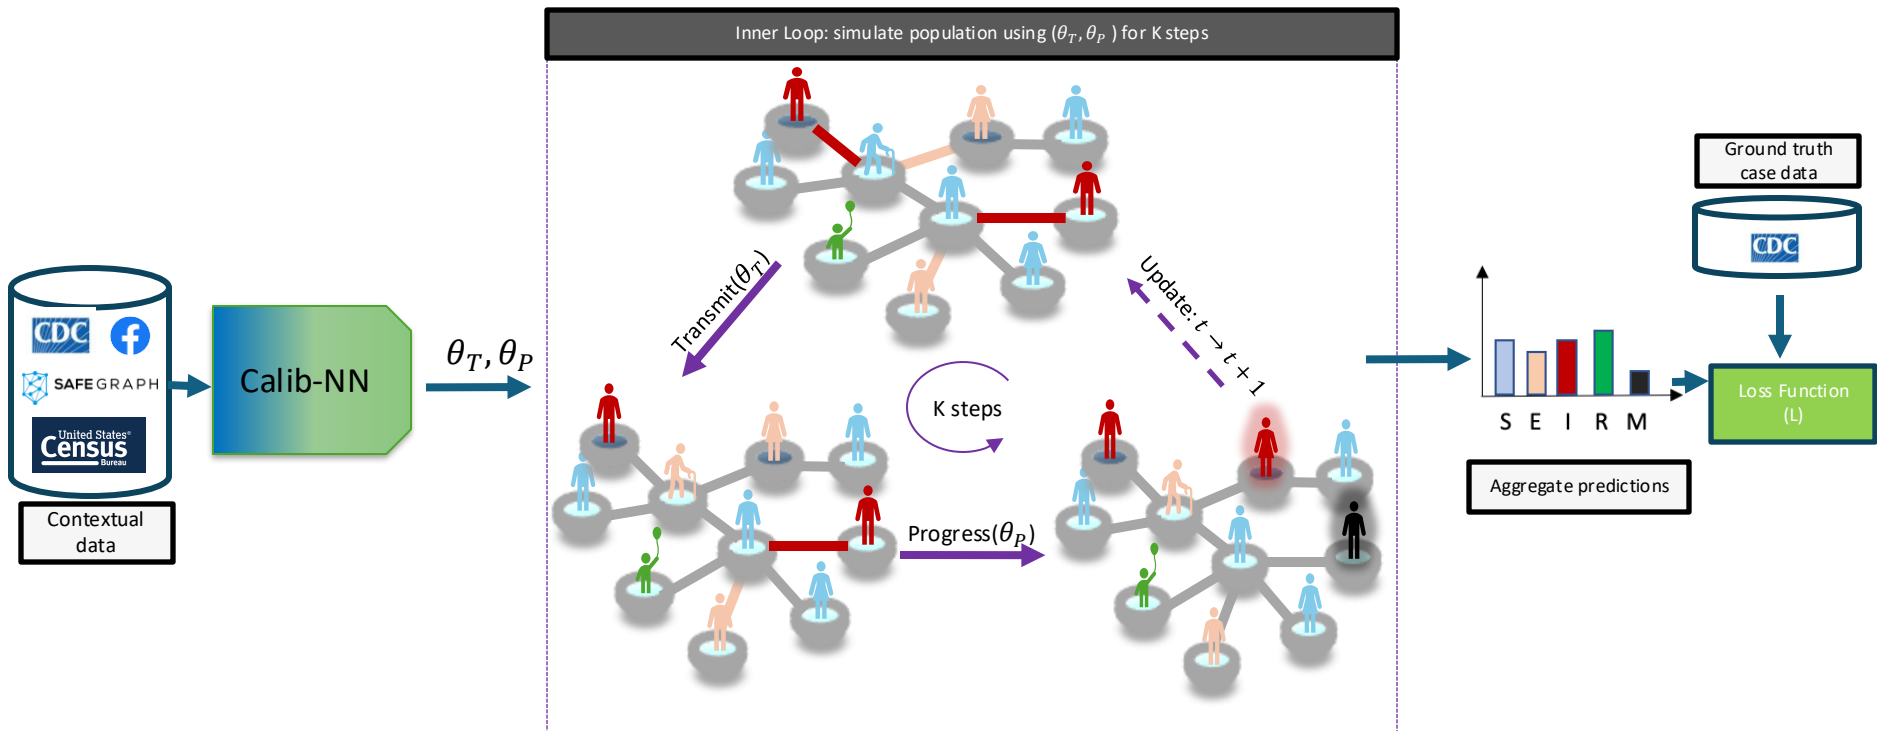


$$\theta_T = \theta_T - \alpha \frac{\partial \mathcal{L}(\hat{y}, y)}{\partial \theta_T}$$

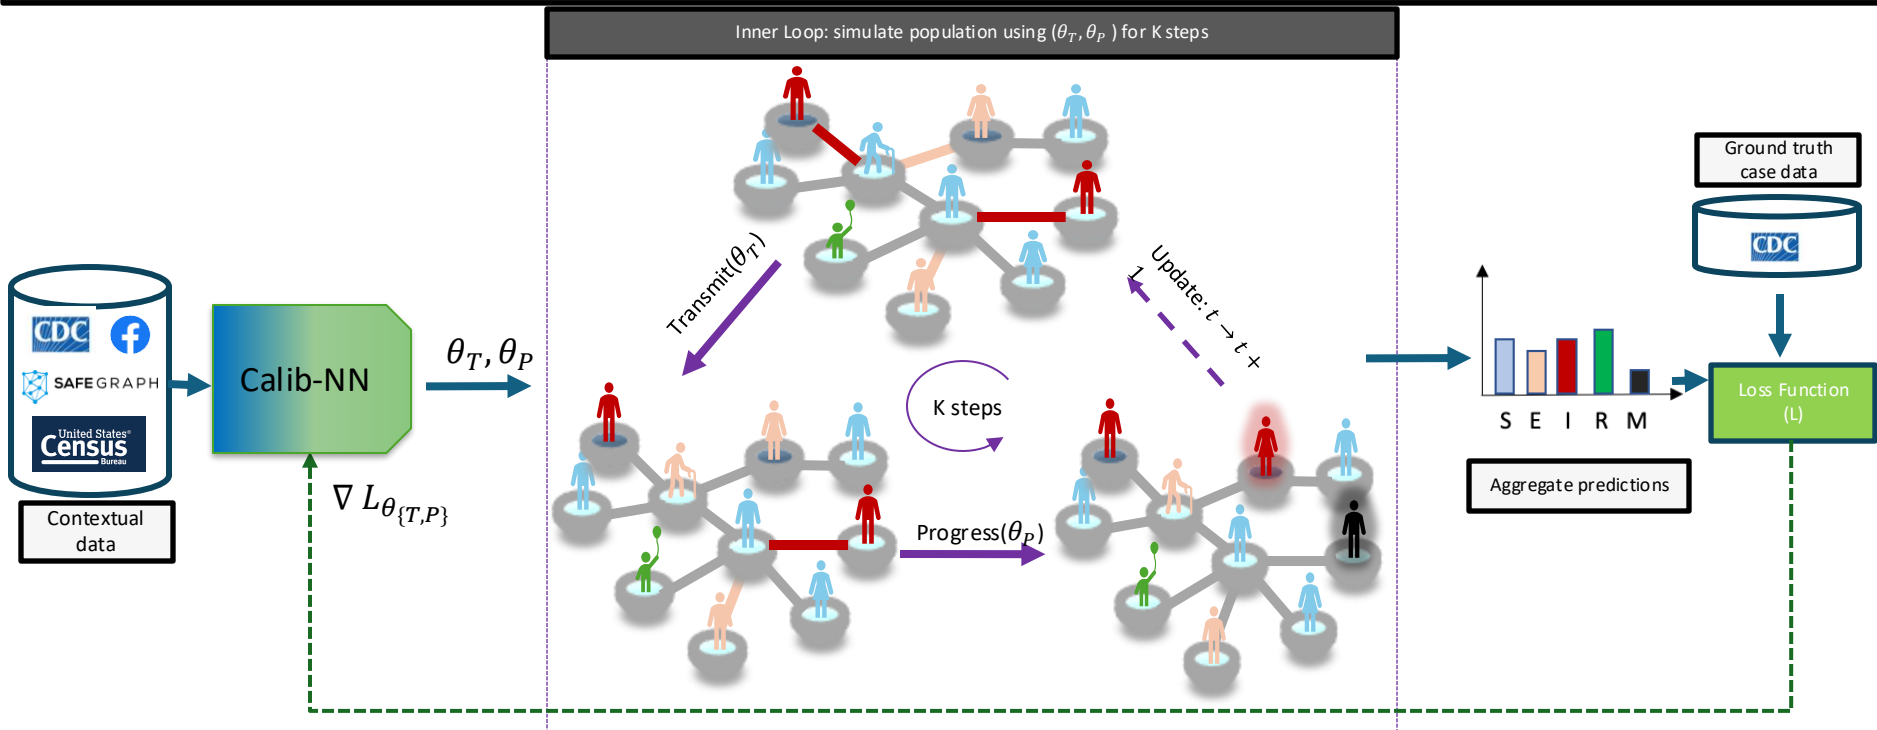
$$\theta_P = \theta_P - \alpha \frac{\partial \mathcal{L}(\hat{y}, y)}{\partial \theta_P}$$

Mode 1: Calibrate **parameters** with gradient descent (c-GRADABM)






Outer Loop: Calib-NN predict infection parameters (θ_T, θ_P) for population simulation and is optimized using end-to-end gradient flow



$$\phi = \phi - \alpha \frac{\partial \mathcal{L}(\hat{y}, y; (\theta_T^t, \theta_P^t))}{\partial \phi},$$

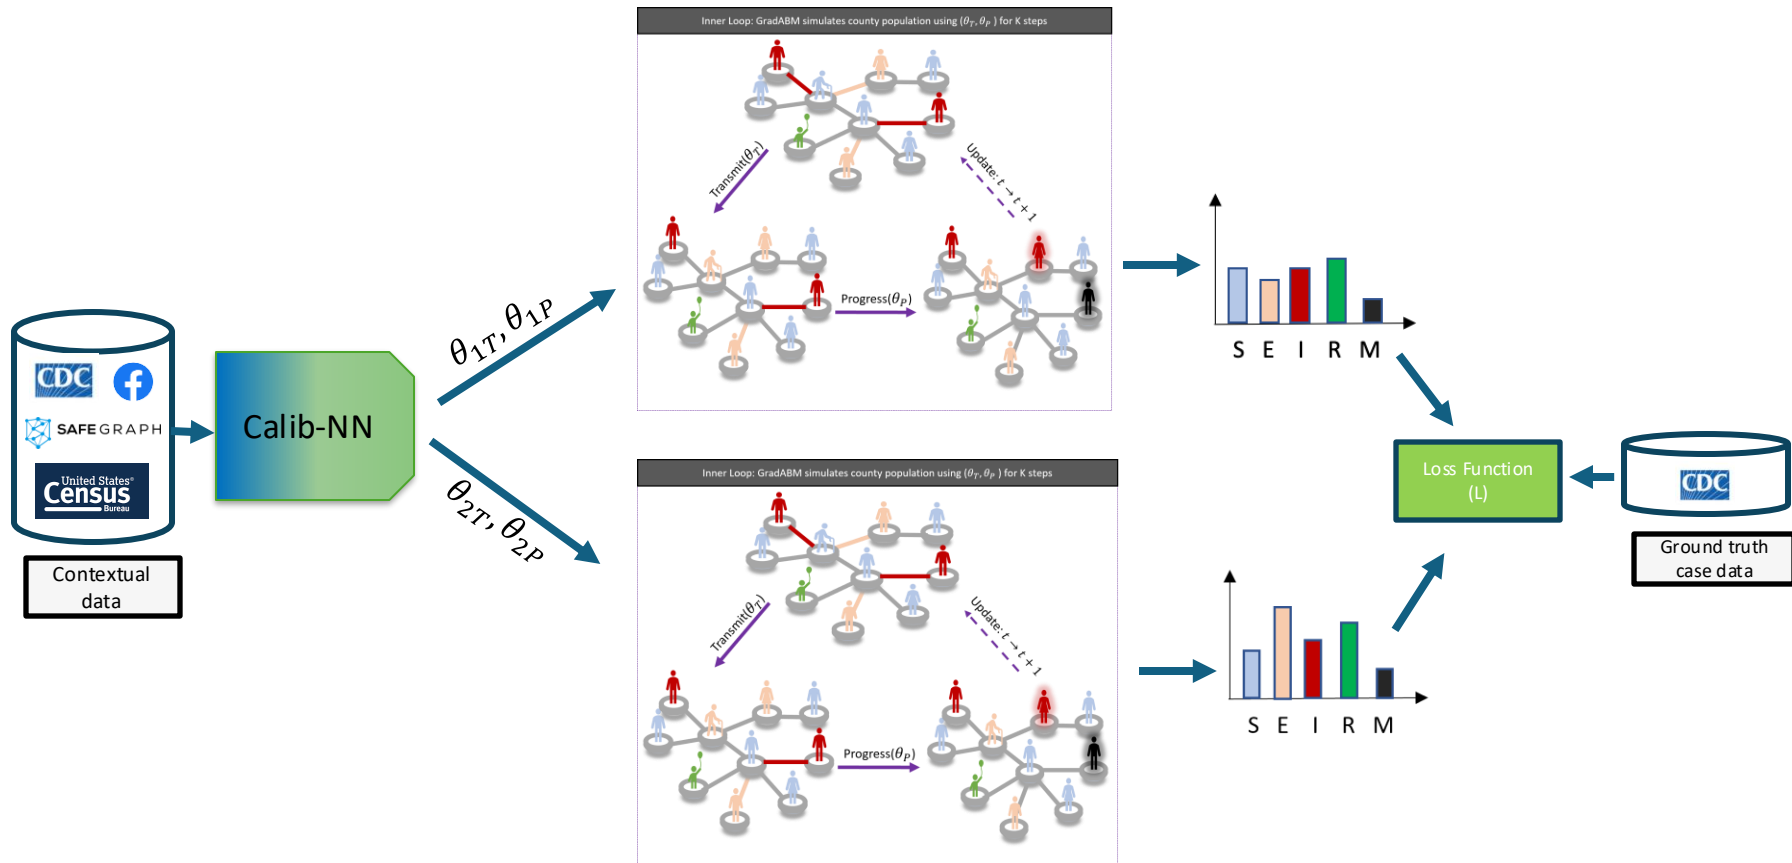
Outer Loop: Calib-NN predict infection parameters (θ_T, θ_P) for population simulation and is optimized using end-to-end gradient flow




Mode 2: Calibrate generator function with gradient descent (dc-GRADABM)

$$\phi = \phi - \alpha \frac{\partial \mathcal{L}(\hat{y}, y; (\theta_T^t, \theta_P^t))}{\partial \phi},$$

Multi-task learning over simulation environments, reduce overfitting



Estimate posteriors over simulation parameters



github.com/AgentTorch/AgentTorch

`pip install agent-torch`





Dynamic interactions

Multi-scale
Stochastic

Differentiability

Compose with ODEs

```
from chirho.dynamical.handlers.solver import TorchDiffEq
from chirho.dynamical.ops import simulate

simulation = envs.create(
    model=epidemiology,
    populations=NYC,
    archetype={'immune_dynamics':
               archetype.ode(chiro_ode, eval_times)
    })

simulation.execute()
```

$$\mathbf{z}_i(t+1) = f\left(\mathbf{z}_i(t), \bigoplus_{j \in \mathcal{N}_i(t)} M_{ij}(t), \boldsymbol{\theta}, \ell(\cdot | \mathbf{z}_i(t))\right)$$

Q3: what about the data?

coarse-grained, noisy and stale data
limited granularity from privacy not scarcity?

$$\mathbf{z}_i(t+1) = f \left(\mathbf{z}_i(t), \bigoplus_{j \in \mathcal{N}_i(t)} M_{ij}(t), \boldsymbol{\theta}, \ell(\cdot | \mathbf{z}_i(t)) \right)$$

- Interaction traces
- Vaccination, disease status

Q3: what about the data?

coarse-grained, noisy and stale data
limited granularity from privacy not scarcity?

$$\mathbf{z}_i(t+1) = f \left(\mathbf{z}_i(t), \bigoplus_{j \in \mathcal{N}_i(t)} M_{ij}(t), \boldsymbol{\theta}, \ell(\cdot | \mathbf{z}_i(t)) \right)$$

- Interaction traces
- Vaccination, disease status

decentralize the simulation?

Large Population Models



Expressive agents

Adaptive
Heterogeneous

LLM Archetypes

AAMAS'23; arxiv'24



Dynamic interactions

Multi-scale
Stochastic

Differentiability

BMJ'21; AAMAS'23;
AAMAS'24 (Best paper runner-up)



Data-driven analysis

Decentralized
Sensitive

Secure MPC

ICML-W'22 (Best paper);
AAMAS'24; Nature Medicine'25



Data-driven analysis

Decentralized
Sensitive

Secure MPC

Backprop through the “real-world”

```
from agent_torch.core import decentralize

# map simulation to real protocol
real_sim = decentralize(
    simulation,
    protocol="ble_contact_tracing")

real_sim.deploy()


real_sim.sync()
```

Compose with Physical protocols

Estimate simulation gradients via MPC

aside: what is secure multi-party computation?

Alice




2

Bob



3


Charlie



5

n=11

Alice




2

Bob



3


Charlie



5

n=11

Alice



2

7 + 5 + 1


Bob



3

2 + 0 + 1

Charlie




5

3 + 1 + 1

n=11

Alice



2

7 + 5 + 1

7 + 2 + 3

Bob




3

2 + 0 + 1

5 + 0 + 1

Charlie




5

3 + 1 + 1

1 + 1 + 1

n=11

Alice



2

7 + 5 + 1

7 + 2 + 3

1

Bob




3

2 + 0 + 1

5 + 0 + 1

6

Charlie



5

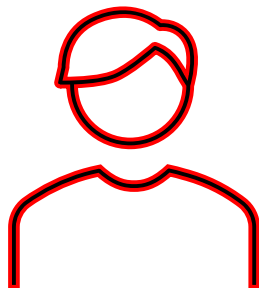
3 + 1 + 1

1 + 1 + 1

3

n=11

Alice



2

$$7 + 5 + 1$$

$$7 + 2 + 3$$

1

$$1 + 6 + 3$$

Bob



3


$$2 + 0 + 1$$

$$5 + 0 + 1$$

6

$$1 + 6 + 3$$

Charlie



5

$$3 + 1 + 1$$


$$1 + 1 + 1$$

3

$$1 + 6 + 3$$

n=11

Alice



2

7 + 5 + 1

7 + 2 + 3

1

answer

10

Bob



3


2 + 0 + 1

5 + 0 + 1

6

10

Charlie



5

3 + 1 + 1


1 + 1 + 1

3

10

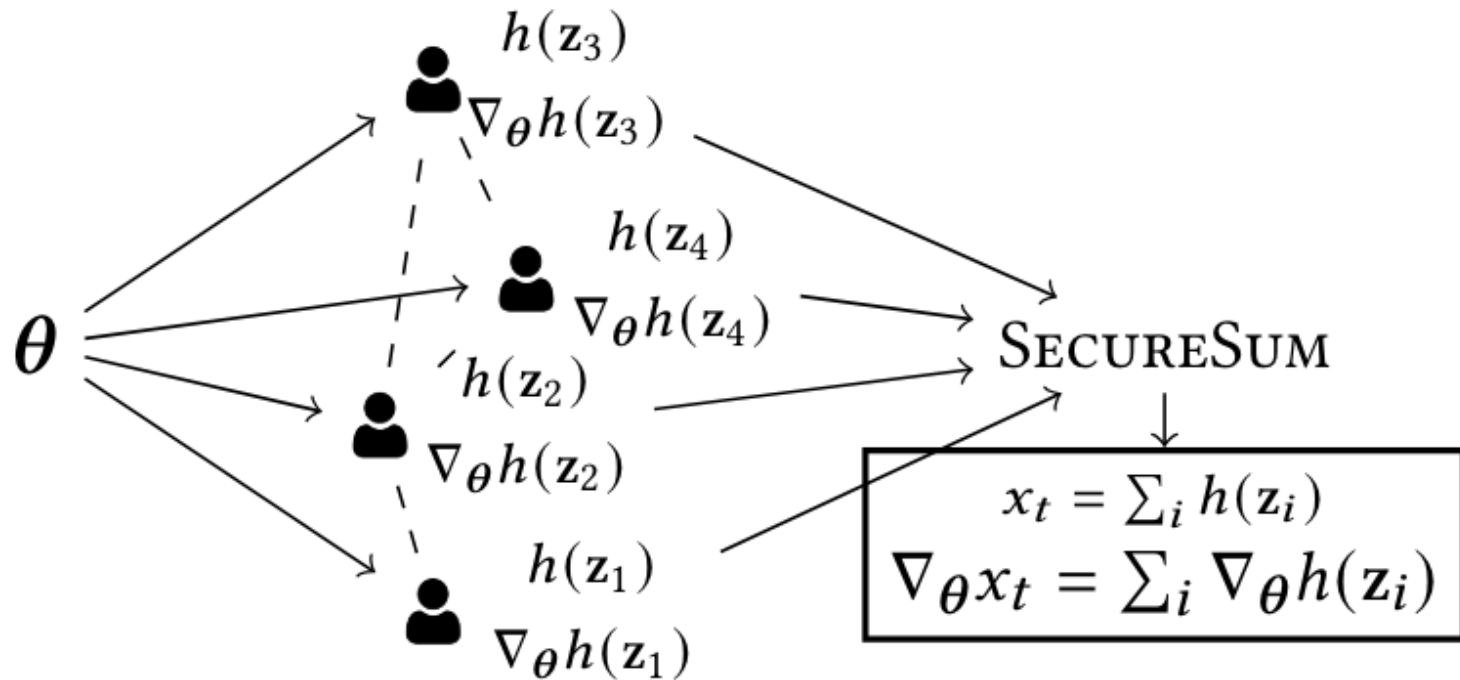
$$\mathbf{z}_i(t+1) = f \left(\mathbf{z}_i(t), \bigoplus_{j \in \mathcal{N}_i(t)} M_{ij}(t), \boldsymbol{\theta}, \ell(\cdot | \mathbf{z}_i(t)) \right)$$

additive secret sharing





securely estimate $z_i(t+1)$ and $dz/d\theta$ on agent device



Aggregate Message and Calibration Gradient



Calibrate disease parameters without leaking an agent's state or interaction trace



Analyze dynamics without leaking individual disease, demographic or geo-location





Data-driven analysis

Decentralized
Sensitive

Secure MPC

Backprop through the “real-world”

```
from agent_torch.core import decentralize

# map simulation to real protocol
real_sim = decentralize(
    simulation,
    protocol="ble_contact_tracing")

real_sim.deploy()

real_sim.sync()
```

Compose with Physical protocols

Estimate simulation gradients via MPC

Large Population Models

lpm.media.mit.edu/research.pdf




Expressive agents

Adaptive
Heterogeneous

LLM Archetypes

AAMAS'23; arxiv'24



Dynamic interactions

Multi-scale
Stochastic

Differentiability

BMJ'21; AAMAS'23;
AAMAS'24 (Best paper runner-up)



Data-driven analysis

Decentralized
Sensitive

Secure MPC

ICML-W'22 (Best paper);
AAMAS'24; Nature Medicine'25

Conclusion: Power of Large Population Models

- Scale and Speed: LPMs simulate millions of agents in seconds on standard hardware, enabling unprecedented insights into complex systems.
- Expressive Agents: By leveraging LLMs as behavioral archetypes, LPMs capture nuanced individual behaviors at scale.
- Closing Sim2Real Gap: Secure multi-party computation allows LPMs to incorporate real-world data without compromising privacy.
- Diverse Applications: From pandemic response to energy adoption and supply chain management, LPMs are already making global impact.
- Open-Source Accessibility: AgentTorch democratizes LPM technology, making it available to researchers and policymakers worldwide.

github.com/AgentTorch/AgentTorch

lpm.media.mit.edu/join

github.com/AgentTorch/AgentTorch

`pip install agent-torch`

